

# Physical Pragmatics (observer\_0)

## Preprocessing

### Re-Label Participant Data

```
# Read in the participant data.
human_0 = read_csv(file.path(human_path, "raw_data.csv"))

# Extract the participant data from the mentalistic condition.
human_1 = human_0 %>%
  filter(target_1!=-1) %>%
  select(workerid, layout, target_0, target_1) %>%
  rename(reward_0=target_0, cooperation=target_1)

# Extract the participant data from the non-mentalistic condition.
human_2 = human_0 %>%
  filter(target_1==1) %>%
  select(workerid, layout, target_0) %>%
  rename(reward_1=target_0) %>%
  mutate(workerid=workerid-length(unique(workerid)))

# Merge the two data partitions and extract the layout information.
human_3 = human_1 %>%
  left_join(human_2)

# Extract the layout information and re-label the natural costs.
human_4 = human_3 %>%
  mutate(layout=substr(layout, 2, 12)) %>%
  separate(layout, into=c("natural_cost", "enforcer_action"), sep="_") %>%
  mutate(natural_cost=factor(natural_cost,
                             levels=c("[5 5]", "[5 7]", "[5 9]",
                                       "[7 5]", "[7 7]", "[7 9]",
                                       "[9 5]", "[9 7]", "[9 9]"),
                             labels=c("[1.25 1.25]", "[1.25 1.75]",
                                       "[1.25 2.25]", "[1.75 1.25]",
                                       "[1.75 1.75]", "[1.75 2.25]",
                                       "[2.25 1.25]", "[2.25 1.75]",
                                       "[2.25 2.25]"))

# Write the preprocessed participant data.
write_csv(human_4, file.path(human_path, "data.csv"))

# Read in the participant age information.
age = read_csv(file.path(human_path, "subject_information.csv")) %>%
  select(workerid, age) %>%
  mutate(age=as.numeric(substr(age, 2, 3)))
```

## Compute Mean Participant Judgments

```
# Read in the preprocessed participant data.
human_4 = read_csv(file.path(human_path, "data.csv"))

# Define the bootstrap functions
compute_mean = function(data, indices) {
  return(mean(data[indices]))
}

compute_bootstrap = function(data) {
  simulations = boot(data=data,
                     statistic=compute_mean,
                     R=10000)

  return(boot.ci(simulations, type="bca")$bca)
}

# Compute the bootstrapped 95% CIs for each dependent measure.
set.seed(seed)
ci = data.frame()
for (nc in unique(human_4$natural_cost)) {
  for (ea in unique(human_4$enforcer_action)) {
    # Filter the relevant data using the current natural cost and enforcer
    # action.
    human_5 = human_4 %>%
      filter(natural_cost==nc, enforcer_action==ea)

    # Compute the bootstrapped 95% CI for each dependent measure.
    reward_0_bootstrap = compute_bootstrap(human_5$reward_0)
    cooperation_bootstrap = compute_bootstrap(human_5$cooperation)
    reward_1_bootstrap = compute_bootstrap(human_5$reward_1)
    ci = rbind(ci, data.frame(natural_cost=nc,
                             enforcer_action=ea,
                             lower_ci_reward_0=reward_0_bootstrap[4],
                             upper_ci_reward_0=reward_0_bootstrap[5],
                             lower_ci_cooperation=cooperation_bootstrap[4],
                             upper_ci_cooperation=cooperation_bootstrap[5],
                             lower_ci_reward_1=reward_1_bootstrap[4],
                             upper_ci_reward_1=reward_1_bootstrap[5]))
  }
}

# Read in the model data.
model_0 = read_csv(file.path(model_path, "predictions.csv"))

# Compute the mean participant judgments for each dependent measure and merge
# the model data.
data_0 = human_4 %>%
  group_by(natural_cost, enforcer_action) %>%
  summarize(human_reward_0=mean(reward_0), human_cooperation=mean(cooperation),
            human_reward_1=mean(reward_1)) %>%
  left_join(ci) %>%
  left_join(select(model_0, -rationality, -enforcer_reward))
```

## Apply Min-Max Scaling

```
# Extract the reward inferences from the mentalistic condition.
mentalistic_rewards = data_0 %>%
  select(natural_cost, enforcer_action, model_reward_0, human_reward_0,
         lower_ci_reward_0, upper_ci_reward_0) %>%
  rename(model=model_reward_0, participants=human_reward_0,
         lower=lower_ci_reward_0, upper=upper_ci_reward_0) %>%
  mutate(type="mentalistic_desires")

# Extract the reward inferences from the non-mentalistic condition.
nonmentalistic_rewards = data_0 %>%
  select(natural_cost, enforcer_action, model_reward_1, human_reward_1,
         lower_ci_reward_1, upper_ci_reward_1) %>%
  rename(model=model_reward_1, participants=human_reward_1,
         lower=lower_ci_reward_1, upper=upper_ci_reward_1) %>%
  mutate(type="non-mentalistic_desires")

# Merge the reward inferences from both conditions.
reward_inferences = mentalistic_rewards %>%
  rbind(nonmentalistic_rewards)

# Apply the min-max scaling to the reward inferences.
model_min = min(reward_inferences$model)
model_max = max(reward_inferences$model) - model_min
human_min = min(reward_inferences$participants)
human_max <- max(reward_inferences$participants) - human_min
data_1 = reward_inferences %>%
  mutate(model=(model-model_min)/model_max,
         participants=(participants-human_min)/human_max,
         lower=(lower-human_min)/human_max,
         upper=(upper-human_min)/human_max)

# Extract the cooperation inferences from the mentalistic condition.
cooperation_inferences = data_0 %>%
  select(natural_cost, enforcer_action, model_cooperation, human_cooperation,
         lower_ci_cooperation, upper_ci_cooperation) %>%
  rename(model=model_cooperation, participants=human_cooperation,
         lower=lower_ci_cooperation, upper=upper_ci_cooperation) %>%
  mutate(type="cooperation")

# Apply the min-max scaling to the cooperation inferences.
model_min = min(cooperation_inferences$model)
model_max = max(cooperation_inferences$model) - model_min
human_min = min(cooperation_inferences$participants)
human_max <- max(cooperation_inferences$participants) - human_min
data_2 = cooperation_inferences %>%
  mutate(model=(model-model_min)/model_max,
         participants=(participants-human_min)/human_max,
         lower=(lower-human_min)/human_max,
         upper=(upper-human_min)/human_max)

# Merge the min-max-scaled data.
data_3 = data_1 %>%
```

```
rbind(data_2)
```


## Analysis of Main Results

Here we plot mean participant judgments ( $N=80$ ,  $M=34.81$  years,  $SD=10.31$  years) against our model's predictions in both conditions jointly.

```
# Plot the model predictions and mean participant judgments in both conditions.
```

```
plot_0 = data_3 %>%  
  ggplot(aes(model, participants)) +  
  geom_point(aes(fill=type), alpha=0.75, pch=21, size=4) +  
  geom_smooth(method="lm", se=FALSE, color="black", linetype="dashed",  
             size=0.5) +  
  theme_classic() +  
  theme(aspect.ratio=1.0) +  
  xlab("Model Predictions") +  
  ylab("Participant Judgments") +  
  scale_fill_manual(name="Inference:",  
                   limits=c("mentalistic_desires", "cooperation",  
                           "non-mentalistic_desires"),  
                   labels=c("Reward (M)", "Cooperation", "Reward (NM)"),  
                   values=c(color_palette[3], color_palette[7],  
                           color_palette[2]))
```

plot\_0



```

# Define the bootstrap functions.
compute_cor = function(data, indices) {
  return(cor(data$model[indices], data$participants[indices],
            method="pearson"))
}

compute_bootstrap = function(data) {
  simulations = boot(data=data,
                    statistic=compute_cor,
                    R=10000)

  return(boot.ci(simulations, type="bca")$bca)
}

# Compute the correlation.
cor_0 = cor(data_3$model, data_3$participants)

# Compute the bootstrapped 95% CI of the correlation.
set.seed(seed)
cor_0_bootstrap = compute_bootstrap(data_3)
cor_0_ci = data.frame(
  lower=cor_0_bootstrap[4],
  upper=cor_0_bootstrap[5]
)

```

Our model predictions yield a correlation of  $r=0.97$  (95% CI: 0.95-0.98) with participant judgments.

## Closer Look at Example Trials

Here we plot the data from a subset of our trials in greater detail.

### *Mentalistic* Condition Examples

We first begin by plotting the data from the *mentalistic* condition.

```

# Extract the data for a given natural cost and clamp the bootstrapped 95% CIs.
data_4 = data_3 %>%
  filter(natural_cost=="[2.25 2.25]") %>%
  gather(source, value, model, participants) %>%
  mutate(lower=ifelse(source=="model", NA, lower),
         lower=ifelse(lower<0, 0, lower),
         upper=ifelse(source=="model", NA, upper),
         upper=ifelse(upper>1, 1, upper))


# Plot the data from the example trials in the mentalistic condition.
plot_1 = data_4 %>%
  filter(type %in% c("mentalistic_desires", "cooperation")) %>%
  ggplot(aes(x=type, y=value, fill=source)) +
  geom_bar(stat="identity", position=position_dodge(), color="black") +
  geom_errorbar(aes(ymin=lower, ymax=upper), position=position_dodge(0.9),
               width=0.3) +
  facet_wrap(~factor(enforcer_action,
                    levels=c("[1 0]", "[2 0]", "[3 0]"),
                    labels=c("1 Boulder", "2 Boulders", "3 Boulders")) +

```

```

theme_classic() +
theme(aspect.ratio=1.0) +
scale_x_discrete(name="Inference",
                 limits=c("mentalistic_desires", "cooperation"),
                 labels=c("Reward (M)", "Cooperation")) +
scale_y_continuous(name="Inferred Value",
                  expand=c(0, 0),
                  limits=c(0, 1.05)) +
scale_fill_manual(name="Source:",
                 limits=c("model", "participants"),
                 labels=c("Model", "Participants"),
                 values=c(color_palette[8], color_palette[3]))
plot_1

```



## Non-Mentalistic Condition Examples

Now we plot the data from the *non-mentalistic* condition.

```


# Plot the data from the example trials in the non-mentalistic condition.
plot_2 = data_4 %>%
  filter(type %in% c("non-mentalistic_desires")) %>%
  ggplot(aes(x=type, y=value, fill=source)) +
  geom_bar(stat="identity", position=position_dodge(), width=0.5,
          color="black") +
  geom_errorbar(aes(ymin=lower, ymax=upper), position=position_dodge(0.5),
               width=0.15) +

```

```

facet_wrap(~factor(enforcer_action,
                    levels=c("[1 0]", "[2 0]", "[3 0]"),
                    labels=c("1 Boulder", "2 Boulders", "3 Boulders")) +
theme_classic() +
theme(aspect.ratio=1.0) +
scale_x_discrete(name="Inference",
                 limits=c("non-mentalistic_desires"),
                 labels=c("Reward (NM)")) +
scale_y_continuous(name="Inferred Value",
                   expand=c(0, 0),
                   limits=c(0, 1.05)) +
scale_fill_manual(name="Source:",
                  limits=c("model", "participants"),
                  labels=c("Model", "Participants"),
                  values=c(color_palette[8], color_palette[3]))
plot_2

```



## Decider Cost Lesion

Here we evaluate mean participant judgments ( $N=80$ ,  $M=34.81$  years,  $SD=10.31$  years) against our model's performance when deciders no longer have an understanding of costs—that is, they navigate the world with their rewards as their only consideration.

## Preprocessing

### Apply Min-Max Scaling

```
# Read in the model predictions of the "decider cost lesion" model.
model_1 = read_csv(file.path(model_path,
                             "decider_cost_lesion/predictions.csv"))

# Compute the mean participant judgments for each dependent measure and merge
# the model data.
data_5 = human_4 %>%
  group_by(natural_cost, enforcer_action) %>%
  summarize(human_reward_0=mean(reward_0), human_cooperation=mean(cooperation),
            human_reward_1=mean(reward_1)) %>%
  left_join(ci) %>%
  left_join(select(model_1, -rationality, -enforcer_reward))

# Extract the reward inferences from the mentalistic condition.
mentalistic_rewards = data_5 %>%
  select(natural_cost, enforcer_action, model_reward_0, human_reward_0,
         lower_ci_reward_0, upper_ci_reward_0) %>%
  rename(model=model_reward_0, participants=human_reward_0,
         lower=lower_ci_reward_0, upper=upper_ci_reward_0) %>%
  mutate(type="mentalistic_desires")

# Extract the reward inferences from the non-mentalistic condition.
nonmentalistic_rewards = data_5 %>%
  select(natural_cost, enforcer_action, model_reward_1, human_reward_1,
         lower_ci_reward_1, upper_ci_reward_1) %>%
  rename(model=model_reward_1, participants=human_reward_1,
         lower=lower_ci_reward_1, upper=upper_ci_reward_1) %>%
  mutate(type="non-mentalistic_desires")

# Merge the reward inferences from both conditions.
reward_inferences = mentalistic_rewards %>%
  rbind(nonmentalistic_rewards)

# Apply the min-max scaling to the reward inferences.
model_min = min(reward_inferences$model)
model_max = max(reward_inferences$model) - model_min
human_min = min(reward_inferences$participants)
human_max <- max(reward_inferences$participants) - human_min
data_6 = reward_inferences %>%
  mutate(model=(model-model_min)/model_max,
         participants=(participants-human_min)/human_max,
         lower=(lower-human_min)/human_max,
         upper=(upper-human_min)/human_max)

# Extract the cooperation inferences from the mentalistic condition.
cooperation_inferences = data_5 %>%
  select(natural_cost, enforcer_action, model_cooperation, human_cooperation,
         lower_ci_cooperation, upper_ci_cooperation) %>%
  rename(model=model_cooperation, participants=human_cooperation,
         lower=lower_ci_cooperation, upper=upper_ci_cooperation) %>%
```



```

mutate(type="cooperation")

# Apply the min-max scaling to the cooperation inferences.
model_min = min(cooperation_inferences$model)
model_max = max(cooperation_inferences$model) - model_min
human_min = min(cooperation_inferences$participants)
human_max <- max(cooperation_inferences$participants) - human_min
data_7 = cooperation_inferences %>%
  mutate(model=(model-model_min)/model_max,
         participants=(participants-human_min)/human_max,
         lower=(lower-human_min)/human_max,
         upper=(upper-human_min)/human_max)

# Merge the min-max-scaled data.
data_8 = data_6 %>%
  rbind(data_7)


```

## Analysis of “Decider Cost Lesion” Model Results

```

# Plot the model predictions and mean participant judgments in both conditions
# for the "decider cost lesion" model.
plot_3 = data_8 %>%
  ggplot(aes(model, participants)) +
  geom_point(aes(fill=type), alpha=0.75, pch=21, size=4) +
  geom_smooth(method="lm", se=FALSE, color="black", linetype="dashed",
             size=0.5) +
  theme_classic() +
  theme(aspect.ratio=1.0) +
  xlab("Model Predictions") +
  ylab("Participant Judgments") +
  scale_fill_manual(name="Inference:",
                   limits=c("mentalistic_desires", "cooperation",
                           "non-mentalistic_desires"),
                   labels=c("Reward (M)", "Cooperation", "Reward (NM)"),
                   values=c(color_palette[3], color_palette[7],
                           color_palette[2]))
plot_3

```



```
# Define the bootstrap functions.
compute_cor = function(data, indices) {
  return(cor(data$model[indices], data$participants[indices],
    method="pearson"))
}

compute_bootstrap = function(data) {
  simulations = boot(data=data,
    statistic=compute_cor,
    R=10000)

  return(boot.ci(simulations, type="bca")$bca)
}

# Compute the correlation.
cor_1 = cor(data_8$model, data_8$participants)

# Compute the bootstrapped 95% CI for the correlation.
set.seed(seed)
cor_1_bootstrap = compute_bootstrap(data_8)
cor_1_ci = data.frame(
  lower=cor_1_bootstrap[4],
  upper=cor_1_bootstrap[5]
)
```

Our “decider cost lesion” model predictions yield a correlation of  $r=0.29$  (95% CI: 0.08-0.47) with participant judgments.

## Enforcer Cost Lesion

Now we evaluate mean participant judgments ( $N=80$ ,  $M=34.81$  years,  $SD=10.31$  years) against our model's performance when enforcers no longer have an understanding of costs—that is, they manipulate the environment with their rewards as their only consideration.

### Preprocessing

#### Apply Min-Max Scaling

```
# Read in the model predictions of the "enforcer cost lesion" model.
model_2 = read_csv(file.path(model_path,
                             "enforcer_cost_lesion/predictions.csv"))

# Compute the mean participant judgments for each dependent measure and merge
# the model data.
data_9 = human_4 %>%
  group_by(natural_cost, enforcer_action) %>%
  summarize(human_reward_0=mean(reward_0), human_cooperation=mean(cooperation),
            human_reward_1=mean(reward_1)) %>%
  left_join(ci) %>%
  left_join(select(model_2, -rationality, -enforcer_reward))

# Extract the reward inferences from the mentalistic condition.
mentalistic_rewards = data_9 %>%
  select(natural_cost, enforcer_action, model_reward_0, human_reward_0,
         lower_ci_reward_0, upper_ci_reward_0) %>%
  rename(model=model_reward_0, participants=human_reward_0,
         lower=lower_ci_reward_0, upper=upper_ci_reward_0) %>%
  mutate(type="mentalistic_desires")

# Extract the reward inferences from the non-mentalistic condition.
nonmentalistic_rewards = data_9 %>%
  select(natural_cost, enforcer_action, model_reward_1, human_reward_1,
         lower_ci_reward_1, upper_ci_reward_1) %>%
  rename(model=model_reward_1, participants=human_reward_1,
         lower=lower_ci_reward_1, upper=upper_ci_reward_1) %>%
  mutate(type="non-mentalistic_desires")

# Merge the reward inferences from both conditions.
reward_inferences = mentalistic_rewards %>%
  rbind(nonmentalistic_rewards)

# Apply the min-max scaling to the reward inferences.
model_min = min(reward_inferences$model)
model_max = max(reward_inferences$model) - model_min
human_min = min(reward_inferences$participants)
human_max <- max(reward_inferences$participants) - human_min
data_10 = reward_inferences %>%
  mutate(model=(model-model_min)/model_max,
         participants=(participants-human_min)/human_max,
         lower=(lower-human_min)/human_max,
         upper=(upper-human_min)/human_max)
```

```

# Extract the cooperation inferences from the mentalistic condition.
cooperation_inferences = data_9 %>%
  select(natural_cost, enforcer_action, model_cooperation, human_cooperation,
         lower_ci_cooperation, upper_ci_cooperation) %>%
  rename(model=model_cooperation, participants=human_cooperation,
         lower=lower_ci_cooperation, upper=upper_ci_cooperation) %>%
  mutate(type="cooperation")

# Apply the min-max scaling to the cooperation inferences.
model_min = min(cooperation_inferences$model)
model_max = max(cooperation_inferences$model) - model_min
human_min = min(cooperation_inferences$participants)
human_max <- max(cooperation_inferences$participants) - human_min
data_11 = cooperation_inferences %>%
  mutate(model=(model-model_min)/model_max,
         participants=(participants-human_min)/human_max,
         lower=(lower-human_min)/human_max,
         upper=(upper-human_min)/human_max)

# Merge the min-max-scaled data.
data_12 = data_10 %>%
  rbind(data_11)


```

## Analysis of Enforcer Cost Lesion Results

```

# Plot the model predictions and mean participant judgments in both conditions
# for the "enforcer cost lesion" model.
plot_4 = data_12 %>%
  ggplot(aes(model, participants)) +
  geom_point(aes(fill=type), alpha=0.75, pch=21, size=4) +
  geom_smooth(method="lm", se=FALSE, color="black", linetype="dashed",
             size=0.5) +
  theme_classic() +
  theme(aspect.ratio=1.0) +
  xlab("Model Predictions") +
  ylab("Participant Judgments") +
  scale_fill_manual(name="Inference:",
                   limits=c("mentalistic_desires", "cooperation",
                           "non-mentalistic_desires"),
                   labels=c("Reward (M)", "Cooperation", "Reward (NM)"),
                   values=c(color_palette[3], color_palette[7],
                           color_palette[2]))
plot_4

```



```
# Define the bootstrap functions.
compute_cor = function(data, indices) {
  return(cor(data$model[indices], data$participants[indices],
    method="pearson"))
}

compute_bootstrap = function(data) {
  simulations = boot(data=data,
    statistic=compute_cor,
    R=10000)

  return(boot.ci(simulations, type="bca")$bca)
}

# Compute the correlation.
cor_2 = cor(data_12$model, data_12$participants)

# Compute the bootstrapped 95% CI for the correlation.
set.seed(seed)
cor_2_bootstrap = compute_bootstrap(data_12)
cor_2_ci = data.frame(
  lower=cor_2_bootstrap[4],
  upper=cor_2_bootstrap[5]
)
```

Our “enforcer cost lesion” model predictions yield a correlation of  $r=0.91$  (95% CI: 0.86-0.94) with participant judgments.

## Full Model vs. Lesioned Models

```
# Define the bootstrap functions.
compute_cor_diff = function(data, indices) {
  cor_full_model = cor(data$full_model[indices],
                        data$participants[indices],
                        method="pearson")
  cor_lesioned_model = cor(data$lesioned_model[indices],
                           data$participants[indices],
                           method="pearson")
  return(cor_full_model-cor_lesioned_model)
}

# Define the bootstrap function to simulate the data.
compute_bootstrap = function(data) {
  simulations = boot(data=data,
                     statistic=compute_cor_diff,
                     R=10000)

  return(boot.ci(simulations, type="bca")$bca)
}

# Merge the full model data with the data from the "decider cost lesion" model.
data_13 = data_3 %>%
  rename(full_model=model) %>%
  left_join(rename(data_8, lesioned_model=model))

# Compute the correlation difference between the full model and the "decider
# cost lesion" model.
cor_diff_0 = cor_0 - cor_1

# Compute the bootstrapped 95% CI for the correlation difference.
set.seed(seed)
cor_diff_0_bootstrap = compute_bootstrap(data_13)
cor_diff_0_ci = data.frame(
  lower=cor_diff_0_bootstrap[4],
  upper=cor_diff_0_bootstrap[5]
)

# Merge the full model data with the data from the "enforcer cost lesion" model.
data_14 = data_3 %>%
  rename(full_model=model) %>%
  left_join(rename(data_12, lesioned_model=model))

# Compute the correlation difference between the full model and the "enforcer
# cost lesion" model.
cor_diff_1 = cor_0 - cor_2

# Compute the bootstrapped 95% CI for the correlation difference.
set.seed(seed)
cor_diff_1_bootstrap = compute_bootstrap(data_14)
cor_diff_1_ci = data.frame(
  lower=cor_diff_1_bootstrap[4],
  upper=cor_diff_1_bootstrap[5]
```

)

When compared against the “decider cost lesion” model, our model has a correlation difference of  $\Delta r=0.68$  (95% CI: 0.49-0.89). When compared against the “enforcer cost lesion” model, our model as a correlation difference of  $\Delta r=0.06$  (95% CI: 0.03-0.1).


## Comparing Reward Inferences

Here we conduct an additional analyses comparing the difference in mean participant reward inferences ( $N=80$ ,  $M=34.81$  years,  $SD=10.31$  years) across conditions.

```
# Filter out the cooperation data, clamp the bootstrapped 95% CIs, and
# compute average rewards per condition and enforcer action.
data_15 = data_3 %>%
  filter(type!="cooperation") %>%
  mutate(lower=ifelse(lower<0, 0, lower),
         upper=ifelse(upper>1, 1, upper)) %>%
  group_by(type, enforcer_action) %>%
  summarize(mean_model=mean(model),
            mean_participants=mean(participants),
            mean_lower=mean(lower),
            mean_upper=mean(upper))

# Plot the reward comparison averaged per condition and enforcer action.
plot_5 = data_15 %>%
  ggplot(aes(x=enforcer_action, y=mean_participants, group=type, color=type)) +
  geom_point(size=3) +
  geom_line(size=2, alpha=0.5) +
  geom_errorbar(aes(ymin=mean_lower, ymax=mean_upper), width=0.15, size=2,
               alpha=0.5) +
  theme_classic() +
  theme(aspect.ratio=1.0) +
  scale_x_discrete(name="Number of Boulders",
                  limits=c("[1 0]", "[2 0]", "[3 0]"),
                  labels=c(1, 2, 3)) +
  ylab("Inferred Reward") +
  scale_colour_manual(name="Inference:",
                     limits=c("mentalistic_desires",
                              "non-mentalistic_desires"),
                     labels=c("Reward (M)", "Reward (NM)"),
                     values=c(color_palette[3], color_palette[2]))

plot_5
```



We first quantify the mean participant reward differences across conditions by computing a  $t$ -test over each boulder count.

```
# Filter the reward inferences across in the mentalistic condition, averaged
# over the natural cost and enforcer action.
```

```
data_16 = data_3 %>%
  filter(type=="mentalistic_desires")
```

```
# Filter the reward inferences across in the non-mentalistic condition,
# averaged over the natural cost and enforcer action.
```

```
data_17 = data_3 %>%
  filter(type=="non-mentalistic_desires")
```

```
# Compute a t-test across conditions for one boulder.
```

```
t.test(filter(data_16, enforcer_action=="[1 0]")$participants,
       filter(data_17, enforcer_action=="[1 0]")$participants,
       alternative="two.sided")
```

```
##
```

```
## Welch Two Sample t-test
```

```
##
```

```
## data: filter(data_16, enforcer_action == "[1 0]")$participants and filter(data_17, enforcer_action == "[1 0]")$participants
```

```
## t = 11.533, df = 15.951, p-value = 3.774e-09
```

```
## alternative hypothesis: true difference in means is not equal to 0
```

```
## 95 percent confidence interval:
```

```
## 0.1984210 0.2878201
```

```
## sample estimates:
```



```
## mean of x mean of y
## 0.31342790 0.07030733

# Compute a t-test across conditions for two boulders.
t.test(filter(data_16, enforcer_action=="[2 0]")$participants,
       filter(data_17, enforcer_action=="[2 0]")$participants,
       alternative="two.sided")

##
## Welch Two Sample t-test
##
## data: filter(data_16, enforcer_action == "[2 0]")$participants and filter(data_17, enforcer_action == "[2 0]")$participants
## t = 10.22, df = 15.153, p-value = 3.408e-08
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## 0.09855046 0.15043299
## sample estimates:
## mean of x mean of y
## 0.6299291 0.5054374

# Compute a t-test across conditions for three boulders.
t.test(filter(data_16, enforcer_action=="[3 0]")$participants,
       filter(data_17, enforcer_action=="[3 0]")$participants,
       alternative="two.sided")

##
## Welch Two Sample t-test
##
## data: filter(data_16, enforcer_action == "[3 0]")$participants and filter(data_17, enforcer_action == "[3 0]")$participants
## t = 2.1595, df = 13.084, p-value = 0.04995
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## 7.938303e-06 6.628048e-02
## sample estimates:
## mean of x mean of y
## 0.9628369 0.9296927
```

Next, we quantify the observed reward differences by computing a linear mixed-effects regression (predicting individual participant judgments as a function of enforcer action and condition).

```
# Filter and preprocess the relevant individual participant judgments.
data_18 = human_4 %>%
  select(-cooperation) %>%
  rename(mentalistic_desires=reward_0, nonmentalistic_desires=reward_1) %>%
  gather(condition, reward, mentalistic_desires, nonmentalistic_desires) %>%
  mutate(workerid=ifelse(condition=="nonmentalistic_desires", workerid+40,
                        workerid),
         enforcer_action=as.numeric(str_sub(enforcer_action, 2, 2)))

# Compute a linear mixed-effects regression predicting individual participant
# judgments as a function of enforcer action and condition.
model_3 = lmer(formula=reward ~ enforcer_action*condition +
              (1+condition|workerid), data=data_18)
summary(model_3)

## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
```

```
## Formula: reward ~ enforcer_action * condition + (1 + condition | workerid)
## Data: data_18
##
## REML criterion at convergence: -1965.9
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -5.9730 -0.5614  0.0625  0.5856  3.9789
##
## Random effects:
##   Groups   Name                Variance Std.Dev. Corr
##   workerid (Intercept)          0.01264  0.1124
##             conditionnonmentalistic_desires 0.03543  0.1882  -0.67
##   Residual                    0.02079  0.1442
## Number of obs: 2160, groups: workerid, 80
##
## Fixed effects:
##
##              Estimate Std. Error
## (Intercept)    3.155e-01  2.123e-02
## enforcer_action    1.908e-01  5.373e-03
## conditionnonmentalistic_desires -2.018e-01  3.292e-02
## enforcer_action:conditionnonmentalistic_desires 6.168e-02  7.599e-03
##
##              df t value Pr(>|t|)
## (Intercept)    7.034e+01  14.861 < 2e-16 ***
## enforcer_action    2.078e+03  35.504 < 2e-16 ***
## conditionnonmentalistic_desires    1.202e+02  -6.132 1.14e-08 ***
## enforcer_action:conditionnonmentalistic_desires    2.078e+03   8.117 8.06e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##      (Intr) enfrc_ cndtn_
## enfrcr_ctn -0.506
## cndtnnnmnt_ -0.645  0.326
## enfrcr_ct:_  0.358 -0.707 -0.462
##
## Print the coefficients separately for readability.
summary(model_3)$coefficients
```

```
##              Estimate Std. Error
## (Intercept)    0.31551852  0.021230612
## enforcer_action    0.19076389  0.005372955
## conditionnonmentalistic_desires -0.20184259  0.032916559
## enforcer_action:conditionnonmentalistic_desires 0.06168056  0.007598506
##
##              df    t value
## (Intercept)    70.33871  14.861489
## enforcer_action    2077.99967  35.504461
## conditionnonmentalistic_desires    120.17362  -6.131947
## enforcer_action:conditionnonmentalistic_desires    2077.99967   8.117458
##
##              Pr(>|t|)
## (Intercept)    2.164824e-23
## enforcer_action    3.228756e-216
## conditionnonmentalistic_desires    1.142711e-08
## enforcer_action:conditionnonmentalistic_desires    8.064194e-16
```

Finally, we compute two linear regressions (each predicting mean participant judgments as a function of enforcer action and the natural costs) to examine the predictive power of the enforcer action versus the natural cost in each condition.

```
# Compute a linear regression predicting mean participant judgments as a
# function of enforcer action and the natural costs for the mentalistic
# condition.
data_19 = data_3 %>%
  filter(type=="mentalistic_desires") %>%
  mutate(enforcer_action=as.numeric(substr(enforcer_action, 2, 2))) %>%
  separate(natural_cost, into=c("natural_cost_0", "natural_cost_1"),
           sep=" ") %>%
  mutate(natural_cost_0=as.numeric(gsub("[", "", natural_cost_0, fixed=TRUE)),
         natural_cost_1=as.numeric(gsub("[", "", natural_cost_1, fixed=TRUE)))
model_4 = glm(participants~enforcer_action+natural_cost_0+natural_cost_1,
              family="gaussian", data=data_19)
summary(model_4)
```

```
##
## Call:
## glm(formula = participants ~ enforcer_action + natural_cost_0 +
##      natural_cost_1, family = "gaussian", data = data_19)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.048684  -0.030150   0.001812   0.020063   0.064129
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -0.032628   0.041024  -0.795    0.435
## enforcer_action  0.324704   0.007596  42.746 <2e-16 ***
## natural_cost_0   0.021844   0.015192   1.438    0.164
## natural_cost_1  -0.011206   0.015192  -0.738    0.468
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.001038625)
##
## Null deviance: 1.924395  on 26  degrees of freedom
## Residual deviance: 0.023888  on 23  degrees of freedom
## AIC: -103.19
##
## Number of Fisher Scoring iterations: 2
```

```
# Print the coefficients separately for readability.
summary(model_4)$coefficients
```

```
##              Estimate Std. Error  t value    Pr(>|t|)
## (Intercept)   -0.03262805 0.041023856 -0.7953434 4.345455e-01
## enforcer_action  0.32470449 0.007596142 42.7459739 2.021440e-23
## natural_cost_0   0.02184397 0.015192284  1.4378333 1.639491e-01
## natural_cost_1  -0.01120567 0.015192284 -0.7375898 4.682227e-01
```

```
# Compute a linear regression predicting mean participant judgments as a
# function of enforcer action and the natural costs for the non-mentalistic
# condition.
```

```

data_20 = data_3 %>%
  filter(type=="non-mentalistic_desires") %>%
  mutate(enforcer_action=as.numeric(substr(enforcer_action, 2, 2))) %>%
  separate(natural_cost, into=c("natural_cost_0", "natural_cost_1"),
           sep=" ") %>%
  mutate(natural_cost_0=as.numeric(gsub("[", "", natural_cost_0, fixed=TRUE)),
         natural_cost_1=as.numeric(gsub("]", "", natural_cost_1, fixed=TRUE)))
model_5 = glm(participants~enforcer_action+natural_cost_0+natural_cost_1,
              family="gaussian", data=data_20)
summary(model_5)

```

```

##
## Call:
## glm(formula = participants ~ enforcer_action + natural_cost_0 +
##      natural_cost_1, family = "gaussian", data = data_20)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.063704  -0.010264   0.003176   0.016592   0.045753
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -0.384299   0.037082 -10.363 3.88e-10 ***
## enforcer_action  0.429693   0.006866  62.580 < 2e-16 ***
## natural_cost_0   0.046478   0.013733   3.384 0.00255 **
## natural_cost_1  -0.031206   0.013733  -2.272 0.03273 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.0008486199)
##
## Null deviance: 3.357065  on 26  degrees of freedom
## Residual deviance: 0.019518  on 23  degrees of freedom
## AIC: -108.65
##
## Number of Fisher Scoring iterations: 2
# Print the coefficients separately for readability.
summary(model_5)$coefficients

```

```

##              Estimate Std. Error  t value    Pr(>|t|)
## (Intercept)   -0.3842986 0.037082051 -10.363468 3.881185e-10
## enforcer_action  0.42969267 0.006866262  62.580294 3.389114e-27
## natural_cost_0   0.04647754 0.013732523   3.384487 2.552658e-03
## natural_cost_1  -0.03120567 0.013732523  -2.272392 3.272659e-02

```